

LLM-Based Analysis of User Comments for HarmonyOS Applications

AGCOMMENTSPY — From Raw Reviews to Actionable Insights

Lin Wenqi Hu Haozhan Wang Chaohe

CS326 Project

May 29, 2026

Outline

- 1 Introduction
- 2 Contributions
- 3 System Architecture
 - Comment Scraping
 - LLM Classification
 - Topic Clustering
 - Web Dashboard
- 4 Results
- 5 Evaluation & Discussion
- 6 Conclusion & Future Work

Motivation

HarmonyOS Ecosystem


- Huawei's distributed OS — smartphones, tablets, wearables, IoT
- AppGallery hosts **500,000+** native HarmonyOS apps
- User comments are a critical feedback channel for developers

Unique Challenges of HarmonyOS

- **Closed ecosystem:** HDC protocol, limited tooling (no adb)
- **Heterogeneous devices:** phones, foldables, watches, in-vehicle
- **Rapid API evolution:** compatibility issues surface in reviews

Existing tools target Google Play / App Store — HarmonyOS needs its own pipeline.

Our Approach at a Glance



Key Insight

Use UI automation to collect reviews + LLM to classify & cluster → transform **unstructured reviews** into **actionable developer insights**.

Contributions

- ① **UI-automation crawler** — reliably scrapes AppGallery comments via `hmdriver2` + HDC
- ② **LLM classification pipeline** — multi-dimensional sentiment analysis + actionable suggestions
- ③ **Two-stage clustering** — intra-bucket grouping + cross-bucket merging
- ④ **Interactive web dashboard** — aggregated results in intuitive visual format
- ⑤ **Curated dataset** — **10,000+** comments across 8 categories on HarmonyOS

System Architecture — Overview

Comment Scraping → LLM Classification → Topic Clustering → Web Dashboard

HDC + hmdriver2 | LLM API (JSON mode) | 2-stage LLM cluster | Python HTTP + Vanilla JS

Technology Stack

- hmdriver2 + HDC — device automation
- lxml + XPath — UI parsing
- OpenAI-compatible API — LLM calls
- Python 3.13 — all backend logic
- Vanilla JS + SVG — frontend

Data Flow

- Raw comments → JSON files
- Batch LLM classification (10/batch)
- Classification → clustering
- REST API → dashboard

Stage 1: Comment Scraping

How it works

- Launch AppGallery via deep-link URL
- **Configurable swipes** to load reviews (max 500)
- Parse UI hierarchy (JSON → XML, XPath)
- Extract: username, rating, text, date, device
- **Auto-expand** truncated long comments
- Save per-app JSON files

Scraping Loop (simplified)

```
1 def scrape_comments(driver, max_swipes=50):  
2     all_comments = []  
3     for i in range(max_swipes):  
4         driver.swipe(...)  
5         layout = driver.dump_hierarchy()  
6         xml = json2xml(layout)  
7         for node in xml.xpath(COMMENT_XPATH):  
8             c = parse_comment(node)  
9             all_comments.append(c)  
0     return all_comments
```

Stage 2: LLM-Based Classification

Sentiment (3 labels)

- Positive
- Negative
- Neutral

7 Core Dimensions

- 1 Functional Experience
- 2 System & Device Compatibility
- 3 Stability & Performance
- 4 UI & Interaction
- 5 Operations & Policies
- 6 Appreciation
- 7 Content Ecosystem

Key Design Decisions

- **JSON mode** — structured output enforcement
- **Batch processing** — 10 comments/batch
- **3-tier error recovery:**
 1. Extract JSON from markdown fences
 2. Heuristic brace extraction
 3. Secondary LLM repair call
- **Actionable Suggestion** for every negative issue
- **Idempotent IDs** — traceable to raw comments

Stage 3: Two-Stage Topic Clustering

Stage 1 — Intra-bucket Clustering

- Within each (dimension, sentiment) bucket, LLM groups items by root cause
- Up to 18 items per batch → canonical issue description + display name

Stage 2 — Cross-bucket Merging

- Identifies semantically duplicate topics across dimension-sentiment buckets
- Merged clusters receive standardized topic names

Robustness Mechanisms

- **Auto-repair:** malformed outputs → repair pass → singletons
- **Deduplication:** normalize text + compare (dimension, sentiment, pain_point) triples

Stage 4: Web Dashboard

Four Tabs

1 Overview

- ▶ Total comments, classified count
- ▶ Donut chart (sentiment)
- ▶ Bar chart (dimensions)

2 Comments

- ▶ Paginated list + full-text search
- ▶ Username, rating, date, device

1 Analysis

- ▶ Dimension-filtered pain points
- ▶ LLM-generated **actionable suggestions**

2 Clusters


- ▶ Topics by dimension $\times$ sentiment
- ▶ Expandable sample members

Tech Stack

Vanilla JS (no framework) + inline SVG charts

Python HTTP server + REST API endpoints

Results — Dashboard Overview



- **Dataset:** 10,000+ comments across 45+ apps in 8 categories
- **Sentiment distribution** visualized as an SVG donut chart
- **Dimension breakdown** shown as a horizontal bar chart
- One-click access to total/classified counts

Figure: Overview — aggregate statistics and charts

Results — Analysis & Clusters

[Screenshot: Analysis tab]
dimension-filtered pain points
with actionable suggestions

Figure: Analysis — per-dimension issues with suggestions

[Screenshot: Clusters tab]
topic clusters by dimension × sentiment

Figure: Clusters — grouped issues with canonical descriptions

Evaluation & Discussion

Classification

- 7 dimensions balance **granularity** with **actionability**
- Every negative issue → concrete Actionable_Suggestion

Clustering

- Scattered comments on the same issue → **one root-cause cluster**
- Auto-repair ensures malformed outputs are never silently dropped

Conclusion

What We Built

AGCOMMENTSPY — end-to-end pipeline that transforms **unstructured AppGallery reviews** into **actionable developer insights** via UI automation + LLM-based NLP.

Key Takeaways

- UI automation compensates for the **lack of public APIs** on HarmonyOS
- LLMs enable **zero-shot** fine-grained classification of user reviews
- Two-stage clustering turns **scattered comments** into **coherent topics**
- Lightweight web dashboard makes results **accessible to developers**

Future Work

- ① **Automated labeling** — human-annotated gold standard for evaluation
- ② **LLM comparison** — benchmark multiple providers on cost-quality
- ③ **Incremental scraping** — only scrape new comments since last run
- ④ **Developer workflow integration** — auto-file findings to GitHub Issues / Jira

Thank You!

Questions & Discussion

AGCOMMENTSPY — CS326 Project